

CSE 6220 INTRODUCTION TO HIGH PERFORMANCE COMPUTING

A PARALLEL ALGORITHM EXAMPLE & MODEL OF PARALLEL COMPUTATION

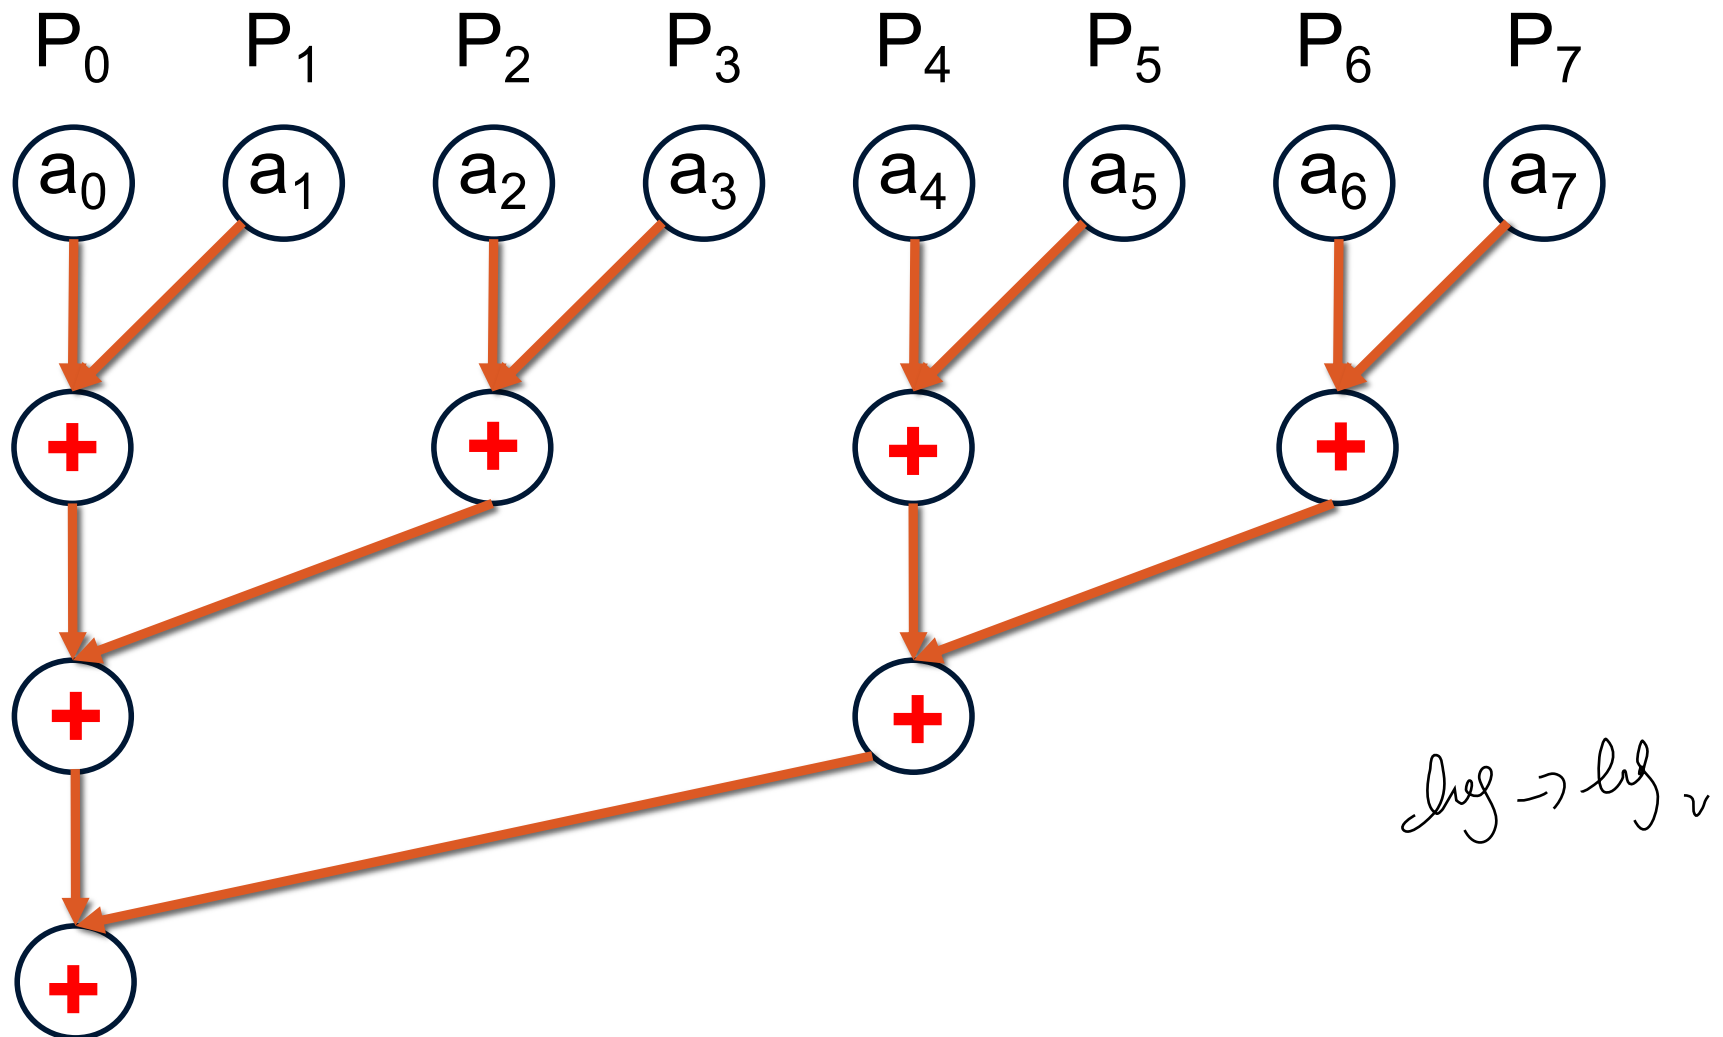
Ümit V. Çatalyürek

School of Computational Science and Engineering
Georgia Institute of Technology

Finding the Sum

- Given n numbers, compute their sum.
- Serial runtime: $T(n, 1) \approx \Theta(n)$

Finding the Sum: A Parallel Algorithm



Finding the Sum

- Given n numbers, compute their sum.

- Serial runtime: $T(n, 1) \approx \underbrace{c_1} \cdot \underbrace{n}_?$

- Parallel runtime: $T(n, n) \approx c_2 \cdot \log n$

- Speedup: $S(n) \approx \frac{c_1 \cdot n}{c_2 \cdot \log n} \approx \frac{c_1}{c_2} \cdot \frac{n}{\log n}$

- c_1 ? c_2

$$S(n) \leq \frac{n}{\log n}$$

Finding the Sum

- Suppose $n = 1,000,000$.

$$S(n) \leq \frac{n}{\log n}$$


$$S(n) \leq \frac{1,000,000}{20}$$

$$S(n) \leq 50,000$$

(5%) $\leftarrow \frac{1}{20} \text{ core}$


Interconnection Network is Important

1. Bus




one at a time

Finding the Sum: A Parallel Algorithm



Interconnection Network is Important

1. Bus



$$T(n, n) = c_2 \cdot \left[\frac{n}{2} + \frac{n}{4} + \dots + 1 \right]$$

$$T(n, n) = c_2 \cdot (n - 1)$$

$$S(n) \approx \frac{c_1 \cdot n}{c_2 \cdot n} = \frac{c_1}{c_2} < 1$$

Worse: c_2 is latency = τ

$$\frac{c_1}{c_2} \approx \frac{1}{10,000}$$

Interconnection Network is Important

2. All-Connected Network

$$T(n, 1) \approx 1 \cdot n$$

$$T(n, n) \approx \tau \cdot \log n$$

$$S(n) \approx \frac{1,000,000}{10,000 \times 20} = 5$$

there is no global memory.

Another Parallel Algorithm

- Use $p \ll n$ processors.
- Assign $\frac{n}{p}$ numbers per processor.
- **Step 1:** Serial — Sum local $\frac{n}{p}$ numbers.
- **Step 2:** Parallel — Add p numbers using p processors.

Interconnection Network is Important

2. All-Connected Network

$$T(n, 1) \approx n$$

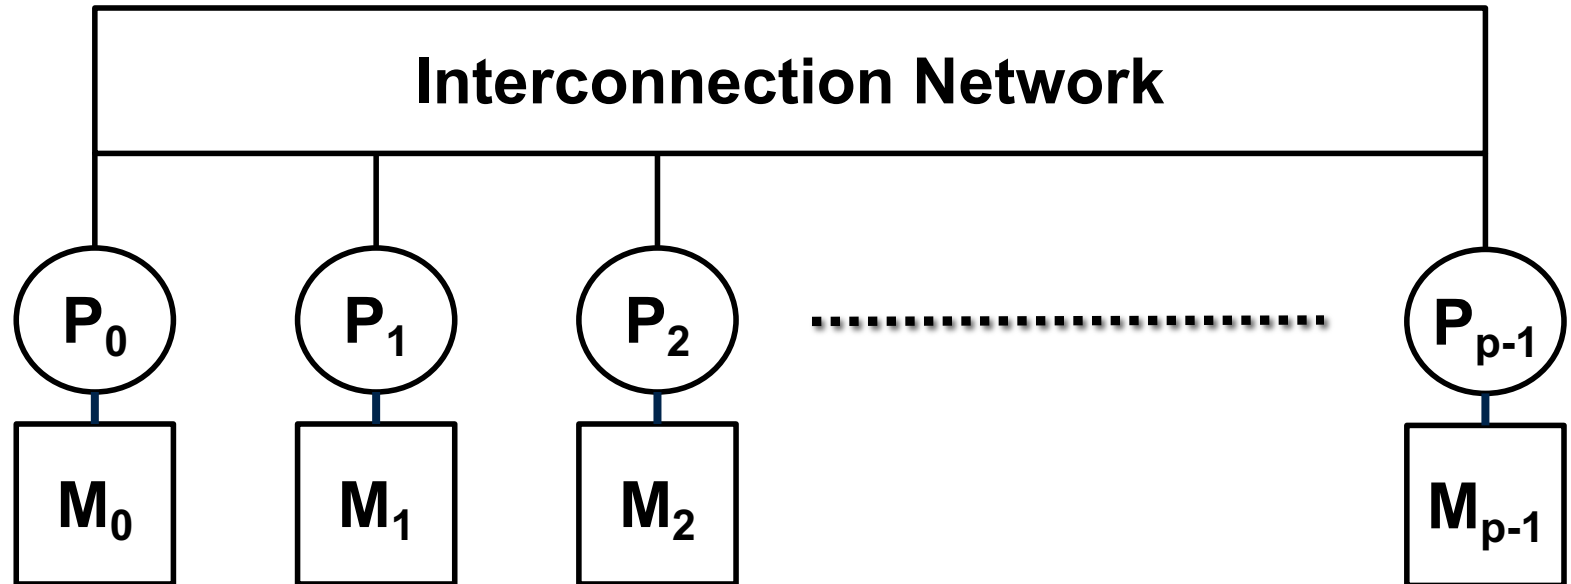
$$T(n, p) \approx \frac{n}{p} + \tau \cdot \log p$$

$$S(p) \approx \frac{n}{\frac{n}{p} + \tau \cdot \log p}$$

$$S(p) \geq \frac{p}{2} \quad \text{if} \quad \tau p \log p \leq n$$

$$\begin{aligned} n &= 1,000,000 \\ p \log p &\leq \frac{n}{\tau} = 100 \\ p &\leq 22 \end{aligned}$$

Model of Parallel Computation



Modeling Communication

- P_i is sending message to P_j of size m
- Transfer time: $t_c = \tau + \mu \cdot m$
 τ (latency) μ (network cost)
- $\frac{1}{\mu}$ = network bandwidth
- 10Gbps = 1.25GBytes/s
- μ (per byte) = $\frac{1}{1.25 \times 10^9} = 0.8 \text{ ns}$

Modeling Communication

- 10Gbps = 1.25GBytes/s
- μ (per byte) = $\frac{1}{1.25 \times 10^9} = 0.8 \text{ ns}$
- μ (per word) = 3.2 ns
- 3GHz \Rightarrow 1 clock cycle = $\frac{1}{3} \text{ ns}$
- $\tau : \mu : 1 \Rightarrow 10^3 - 10^5 : 10 - 20 : 1$

Modeling Communication

- In a parallel communication step
 - Each processor can send at most 1 message
 - Each processor can receive at most 1 message
 - Example suppose $p=8$
 - $0 \rightarrow 5$
 - $1 \rightarrow 0$
 - $2 \rightarrow 1$
 - $3 \rightarrow 7$
 - $4 \rightarrow 2$
 - $5 \rightarrow 4$
 - $6 \rightarrow 6$
 - $7 \rightarrow 3$
- $\rightarrow (5\ 0\ 1\ 7\ 2\ 4\ 6\ 3)$

Permutation Network Model

- Communication that corresponds to any permutation can be realized in parallel.
- $p!$ communication patterns!
- Not everyone has to participate.
- Each message size can be different.
- Not everyone communicates at “exactly” the same time.

Permutation Network Model

- Communication that corresponds to any permutation can be realized in parallel.
- Ideal network?
- Benes $\rightarrow \sim 2 p \log p$ *\leftarrow using serial algorithm, itself.*
- Does not work in practice, requires centralized planner.

Permutation Network Model

1. Shift Permutations:

1. Left shift: $i \rightarrow (i - 1 + p) \bmod p$
2. Right shift: $i \rightarrow (i + 1) \bmod p$

2. Hypercubic Permutations:

- $p = 2^d$ (d is an integer)
- Represent processor ranks with using d -bit binary number
- Example $p = 16$ and $d = 4$
- Consider P_{11} rank = 1011. d positions (0, 1, 2, 3), starting from least significant bit, e.g., position-2 is 0.
- Hypercubic permutation on fixed position j :
 - Two processors communicate if their ranks differ **only** in position j

Hypercubic Permutations

- Example $p = 8, d = 3$

$j = 0$

$0 \leftrightarrow 1$

$2 \leftrightarrow 3$

$4 \leftrightarrow 5$

$6 \leftrightarrow 7$

$j = 1$

$0 \leftrightarrow 2$

$1 \leftrightarrow 3$

$4 \leftrightarrow 6$

$5 \leftrightarrow 7$

$j = 2$

$0 \leftrightarrow 4$

$1 \leftrightarrow 5$

$2 \leftrightarrow 6$

$3 \leftrightarrow 7$

Toy Problem using Hypercubic Permutations

- Given n numbers, compute their sum using p processors.
- Serial runtime: $T(n, 1) = \Theta(n)$
- Parallel runtime: $T(n, p) = \Theta\left(\frac{n}{p} + \log p\right)$

Finding the sum of n numbers

Algorithm (for P_i)

sum \leftarrow add local n/p numbers

for $j=0$ to $d-1$ do

 if ((rank AND 2^j) \neq 0)

 send sum to (rank XOR 2^j)

 else


 receive sum' from (rank XOR 2^j)

 sum = sum + sum'

endfor

if (rank=0)


 print sum



AND, XOR : bitwise operators

*log P
one for.*

Finding the sum of n numbers



Finding the sum of n numbers

Algorithm (for P_i)

sum \leftarrow add local n/p numbers

for $j=0$ to $d-1$ do

 if $((\text{rank AND } 2^j) \neq 0)$

 { send sum to $(\text{rank XOR } 2^j)$; exit; }

 else

 receive sum' from $(\text{rank XOR } 2^j)$

 sum = sum + sum'

endfor

if $(\text{rank}=0)$

 print sum